

Gentry's Fully Homomorphic Encryption Scheme

(over the integers)

Agenda

- Definitions
 - (Fully) Homomorphic Encryption
 - Approximate GCD and SSSP
 - Semantic Security
- History and Applications
- Construction of a Partially Homomorphic Scheme
- Construction of a Fully Homomorphic Encryption Scheme
 - Bootstrappability
 - Tweaking the Initial Scheme
 - Achieving Bootstappability
- Open Questions

Definitions

- Homomorphic Encryption
 - $D(f_i(c_1, c_2, \dots)) = f_i(p_1, p_2, \dots)$
 - f_i is a set of functions
- Fully Homomorphic Encryption
 - f_i *includes every computable function*
- Approximate GCD and SSSP
 - Approximate GCD: Given set of $pq_i + n_i$, determine p
 - SSSP: Given a set of size n and target z , select a subset A of size k where $k \ll n$, $\text{sum}(A) = z$
- Semantic Security
 - Given a plaintext and a ciphertext, it is computationally infeasible to determine if the plaintext encrypts to the ciphertext

History and Applications

- Proposed in 1978, shortly after RSA, by Rivest, Adelman, and Dertouzos.
 - "On Data Banks and Privacy Homomorphisms"
- Partially Homomorphic Schemes: Paillier's, ElGamal, Boneh-Goh-Nissim

RSA



Properties of Homomorphic Encryption

- *Correctness*

- $D(f_i(c_1, c_2, \dots)) = f_i(p_1, p_2, \dots)$
 - Holds true for all f in f_i

- *Compactness*

- The size of the produced ciphertext by *Evaluate* does not depend on the size of the circuit

- *Circuit Privacy*

- The output of *Evaluate* does not reveal anything about the circuit used (except the output itself)

Construction of a Partially Homomorphic Encryption Scheme

Security parameter λ .

Private Key Generation()

Random λ^2 -bit odd integer "k"

Encrypt(k,plaintext)

plaintext is single bit in $\{0,1\}$

m is a random λ -bit integer whose LSB is the plaintext

q is a random λ^5 -bit integer

ciphertext $c = m + kq$

Decryption(k,c)

$(c \bmod k) \bmod 2$

Construction of a Partially Homomorphic Encryption Scheme

- Security of the scheme
 - Approximate GCD
- Efficiency of the scheme
- What makes it partially homomorphic?
 - Turing Machines and circuits
 - Can we do better than a Turing Machine?
- Can we make the scheme a public key encryption scheme?
- What sorts of functions can we compute?
- Can we make the scheme fully homomorphic?
 - If we could find a way to remove noise...

Bootstrappability

We *can* remove noise! Decrypting the ciphertext and re-encrypting the resulting plaintext will do just that!

But what if we don't want private key holder to do *any* work outside the initial encryption and final decryption?

The solution is to have the client decrypt and re-encrypt every time the noise becomes large - homomorphically!

So the question becomes "can the partially homomorphic scheme handle the circuit version of its decryption function?"

For this scheme, not yet.

Tweaking the Initial Scheme

Recall that the decryption function is $(c \bmod k) \bmod 2$.

$$(c \bmod k) \bmod 2 = \text{LSB}(c) \text{ XOR } \text{LSB}(\lceil c/k \rceil)$$

$\lceil c/k \rceil$, then, must be difficult to compute. Is there a way that we can "cheat" so that computing $\lceil c/k \rceil$ will introduce less noise?

What if we introduced a "shortcut" that would allow the server to compute c/k homomorphically?

Tweaking the Initial Scheme

- Key Generation*()
 - $y = \{y_1, y_2, \dots, y_\beta\}$ in $[0, 2)$ such that there exists subset of size α which sums to $1/k$
 - Public key contains y , encryptions of 0, and an encryption of the private key
 - Private key contains k (as before) and the solution s to y
- Encrypt*($k, \text{plaintext}$)
 - Recall that ciphertext $c = m + kq$
 - Compute $z = \{z_1, z_2, \dots, z_i\}$ where $z_i = cy_i$
- Decryption*(k, c)
 - $(c \bmod k) \bmod 2$
 - $= \text{LSB}(k) \text{ XOR } \text{LSB}(\lceil c/k \rceil)$
 - $= \text{LSB}(k) \text{ XOR } \text{LSB}(\lceil \sum s_i z_i \rceil)$

Achieving Bootstrappability

By replacing the two long multiplications by many small additions the encryption scheme becomes bootstrappable for correctly sized parameters:

$$\sum s_i z_i = \sum s_i c y_i = c \sum s_i y_i = c * 1/k = c/k$$

$\sum s_i z_i$ can be obtained by computing a polynomial whose coefficients are the hamming weights of the terms $s_i z_i$.

This computation (decryption) adds noise of $N \alpha \log(\alpha)$ bits. The scheme becomes bootstrappable when the noise is less than $P-4$ bits.

When $\alpha = \lambda / \text{polylog}(\lambda)$, the tweaked scheme becomes feasible.

Open Questions

- Computationally feasible FHE.
 - This scheme requires λ^{10} work.
 - The lattice based scheme requires λ^6 work.
 - Is there a way to lower these requirements?
- FHE without bootstrapping.
 - Bootstrapping, while clever, slows computation.
- FHE over other domains or trapdoor functions.
- FHE with UTM may have very interesting applications.
 - Can we think of other killer cryptographic applications?